
Kernel-bypass techniques for high-speed network packet processing

CS 744

Presenters: Rinku Shah, Priyanka Naik
{rinku, ppnaik}@cse.iitb.ac.in

Course Instructor: Prof. Umesh Bellur

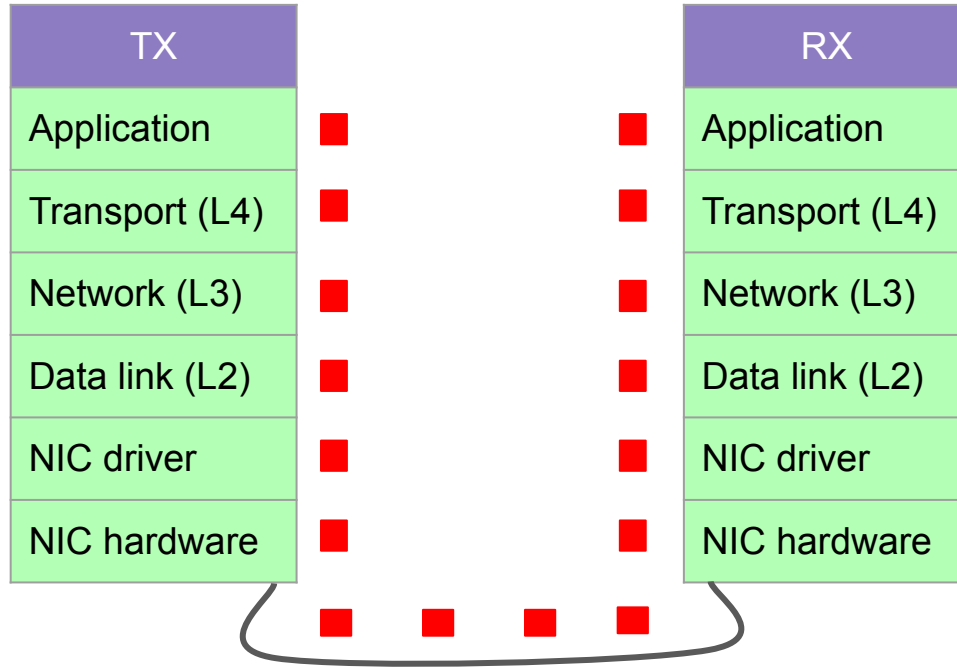


**Department of Computer Science & Engineering
Indian Institute of Technology Bombay**

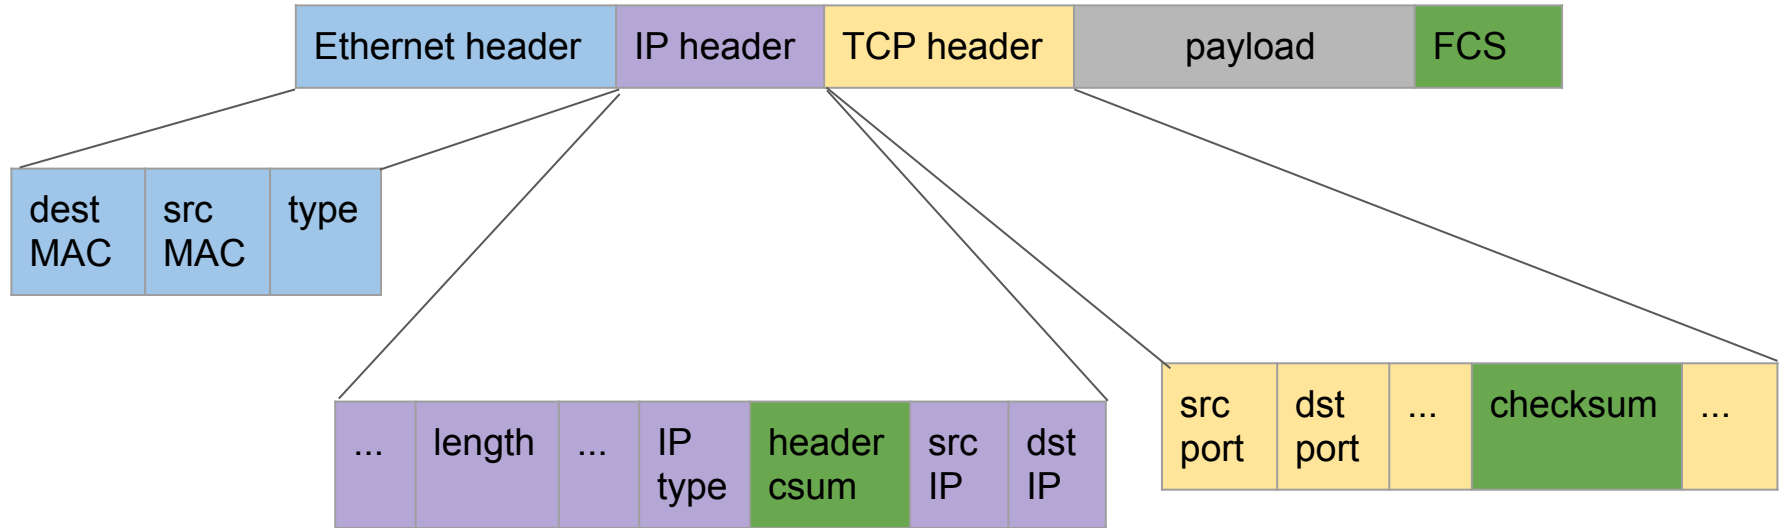
Outline

- The journey of a packet through the Linux network stack
- Need for kernel bypass techniques for packet processing
- Kernel-bypass techniques
 - User-space packet processing
 - **Data Plane Development Kit (DPDK)**
 - **Netmap**
 - User-space network stack
 - **mTCP**
- What's trending?

Typical packet flow



What does a packet contain?

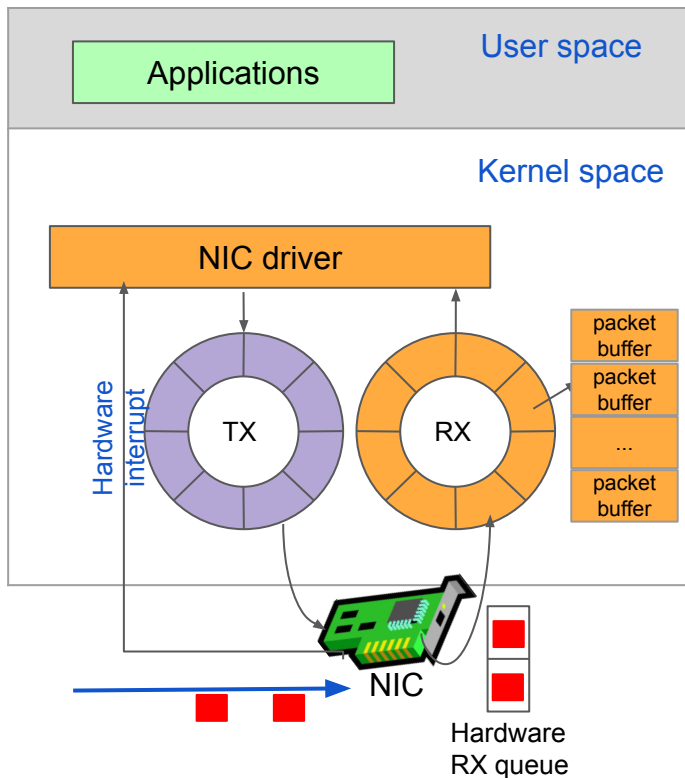


FCS: Frame Check Sequence

Outline

- **The journey of a packet through the Linux network stack**
- Need for kernel bypass techniques for packet processing
- Kernel-bypass techniques
 - User-space packet processing
 - **Data Plane Development Kit (DPDK)**
 - **Netmap**
 - User-space network stack
 - **mTCP**
- What's next??

RX path: Packet arrives at the destination NIC



NIC receives the packet

- Match destination MAC address
- Verify Ethernet checksum (FCS)

Packets accepted at the NIC

- DMA the packet to RX ring buffer
- NIC triggers an interrupt


TX/RX rings

- Circular queue
- Shared between NIC and NIC driver
- Content: Length + packet buffer pointer

Interrupt processing in the linux kernel

- Top-half
 - Minimal processing
- Bottom-half
 - Rest of interrupt processing

Top-half interrupt processing



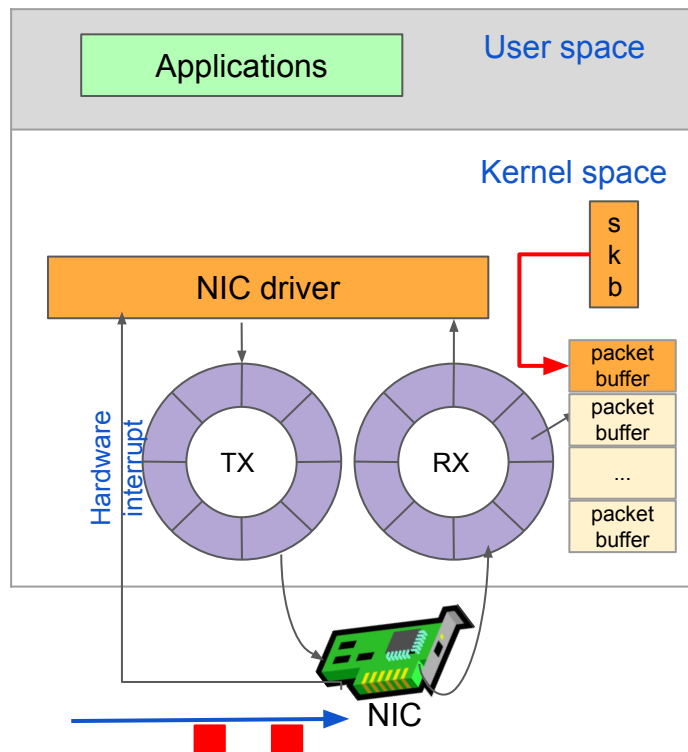
CPU interrupts the process in execution

Switch from user space to kernel space

Top-half interrupt processing

- Lookup IDT (Interrupt Descriptor Table)
- Call corresponding ISR (Interrupt Service Routine)
 - Acknowledge the interrupt
 - Schedule bottom-half processing
- Switch back to user space

Bottom-half processing



CPU initiates the bottom-half when it is free (soft-irq)

Switch from user space to kernel space

Driver dynamically allocates an **sk-buff** (a.k.a., skb)

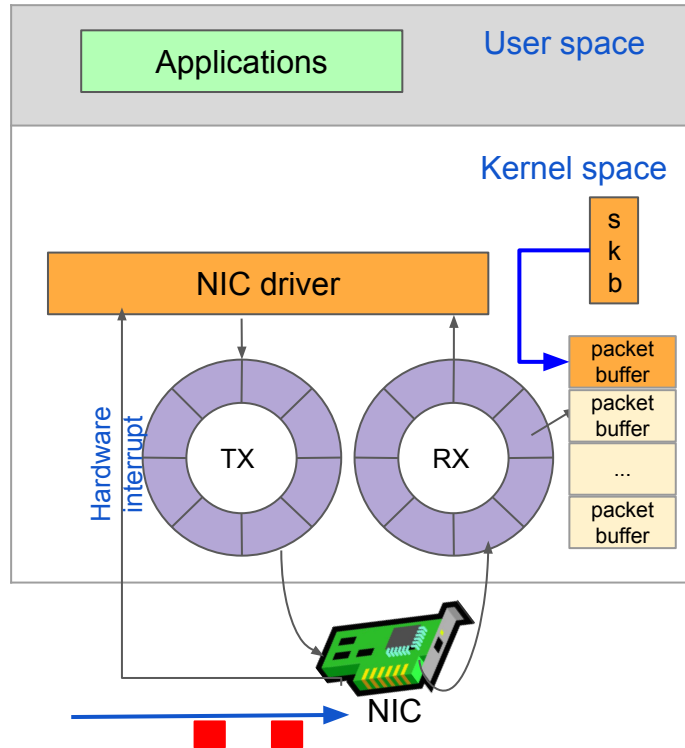


sk-buff ([sk-buff tutorial link](#))

In-memory data structure that contains packet metadata

- Pointers to packet headers and payload
- More packet related information ...

Bottom-half processing



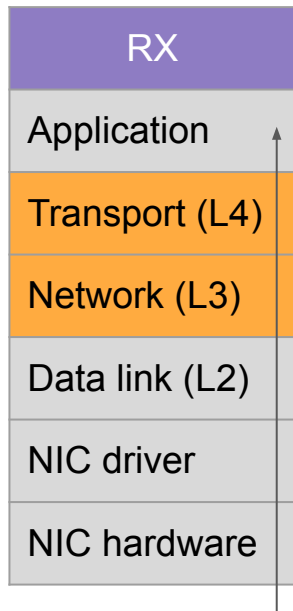
NIC driver processing

For all packets
in buffer

1. Driver dynamically allocates an **skb**
2. Update skb-buff with packet metadata
3. Remove the Ethernet header
4. Pass skb-buff to the network stack

Call L3 protocol handler

L3/L4 processing



Common processing

1. Match destination IP/socket
2. Verify checksum
3. Remove header

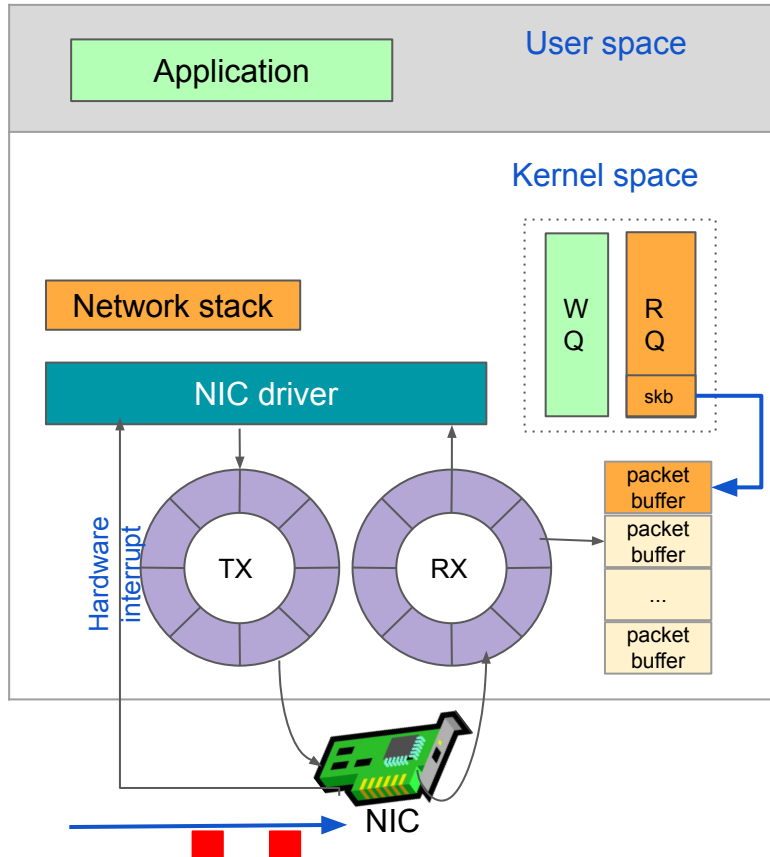


L3-specific processing

1. Route lookup
2. Combine fragmented packets
3. Call L4 protocol handler

L4-specific processing

L3/L4 processing



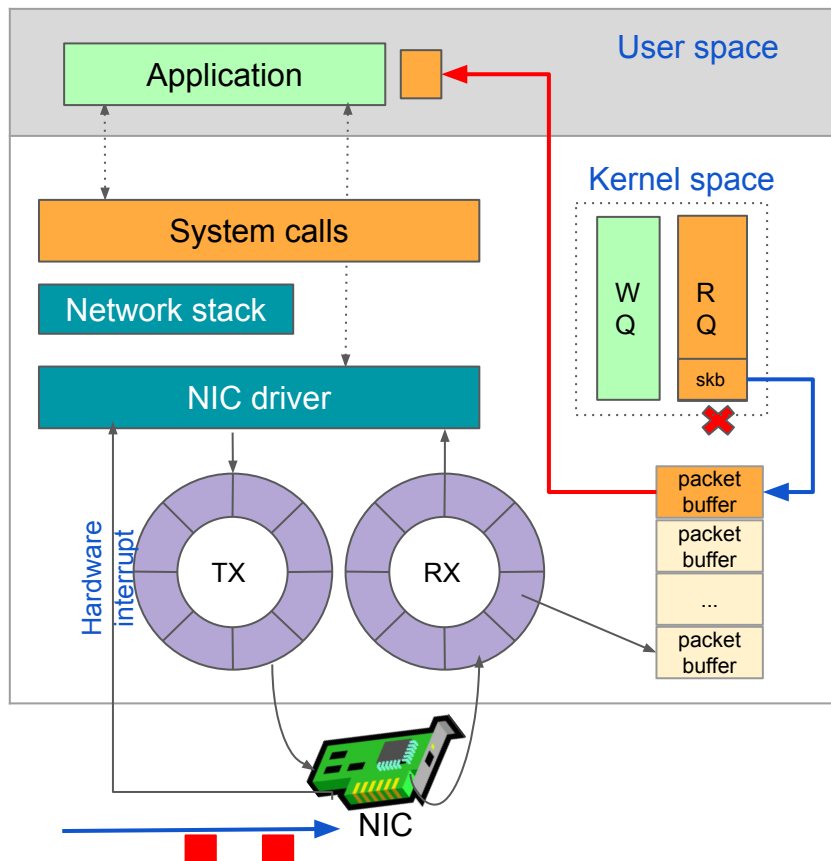
L3-specific processing

1. Route lookup
2. Combine fragmented packets
3. Call L4 protocol handler

L4-specific processing

1. Handle TCP state machine
2. Enqueue to socket read queue
3. Signal the socket

Application processing



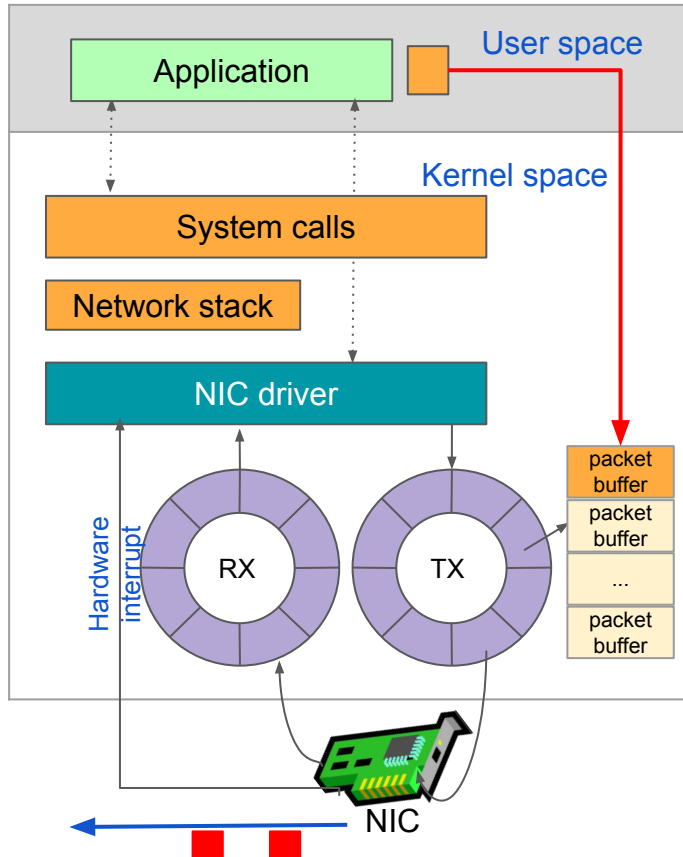
On socket read:

user space to kernel space

- Dequeue packet from socket receive queue (kernel space)
- Copy packet to application buffer (user space)
- Release sk-buff
- Return back to the application

kernel space to user space

Transmit path of an application packet

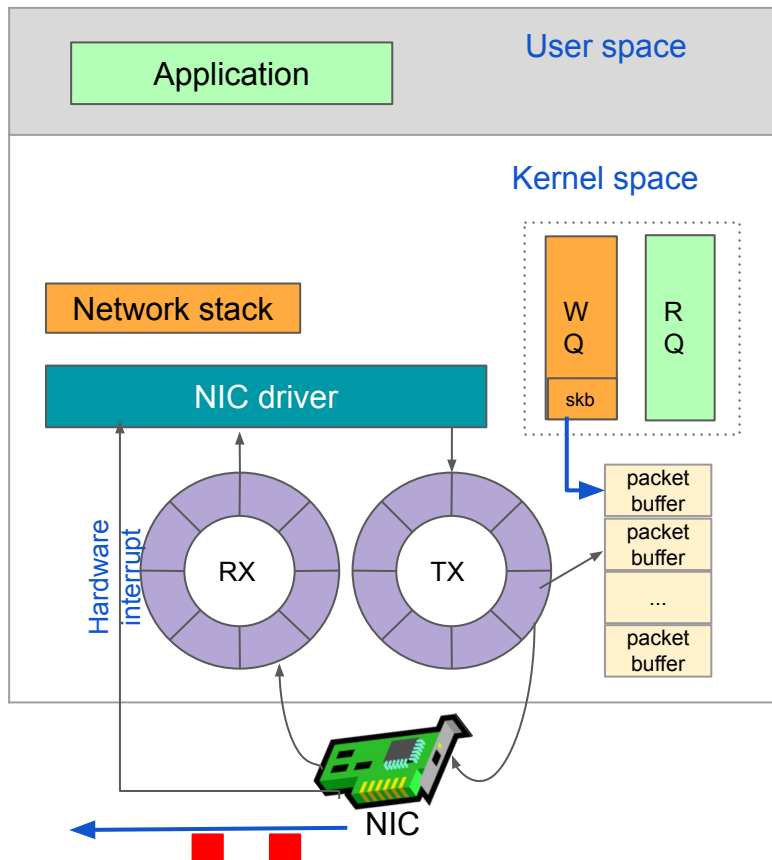


On socket write:

user space to kernel space

- Writes the packet to the kernel buffer
- Calls socket's send function (e.g., sendmsg)

L4/L3 processing



L4-specific processing

1. Allocate sk-buff
2. Enqueue sk-buff to socket write queue
3. Call L3 protocol handler

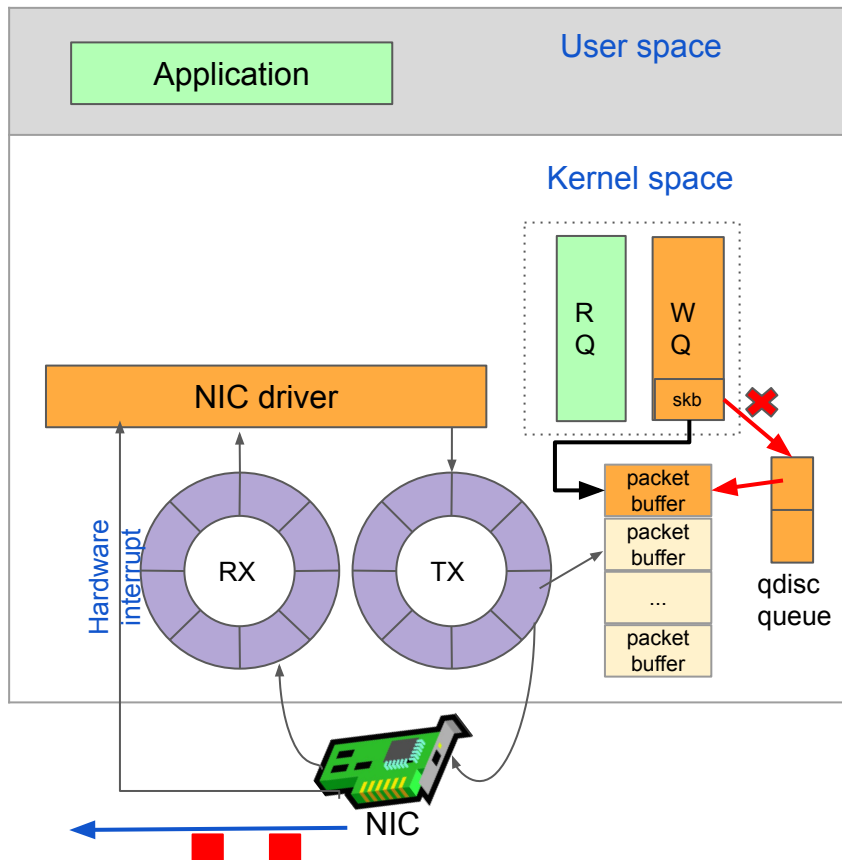
Common processing

1. Build header
2. Add header to packet buffer
3. Update sk-buff

L3-specific processing

1. Fragment, if needed
2. Call L2 protocol handler

L2 processing



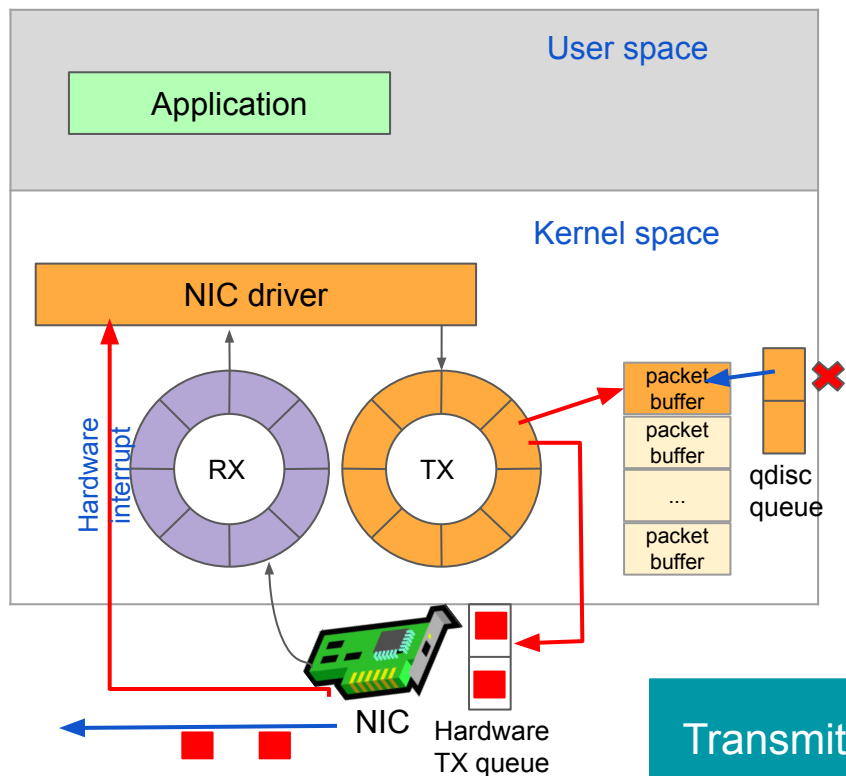
Enqueue packet to **queue discipline** (qdisc)

- Hold packets in a queue
- Apply scheduling policies (e.g. FIFO, priority)

qdisc

- Dequeue sk-buff (if NIC has free buffers)
- Post process sk-buff
 - Calculate IP/TCP checksum
 - ... (tasks that h/w cannot do)
- Call NIC driver's send function

NIC processing



NIC driver

- If hardware transmit queue full
 - Stop qdisc queue
- Otherwise:
 - Map packet data for DMA
 - Tells NIC to send the packet

NIC

- Calculates ethernet frame checksum (FCS)
- Sends packet to the wire
- **Sends an interrupt “Packet is sent” (kernel space to user space)**
- Driver frees the sk-buff; starts the qdisc queue

Transmit and receive packet processing pipeline DONE!!

Packet processing overheads in the kernel

- Too many context switches!!
 - Pollutes CPU cache
- Per-packet interrupt overhead
- Dynamic allocation of sk-buff
- Packet copy between kernel and user space
- Shared data structures

Cannot achieve line-rate for recent high speed NICs!! (40Gbps/100Gbps)

Optimizations to accelerate kernel packet processing

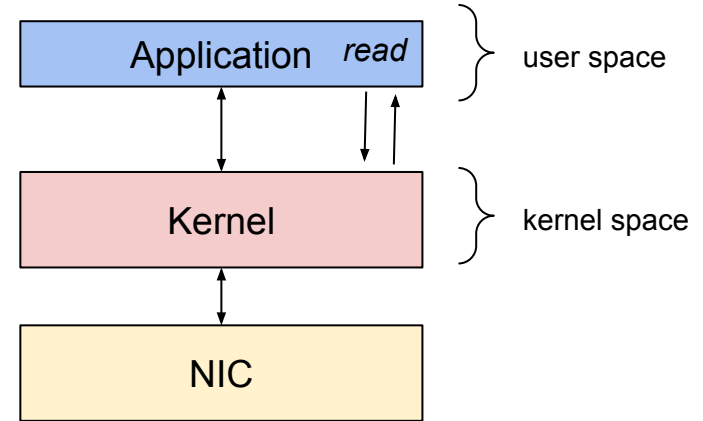
- NAPI (New API) [Reading link](#)
- GRO (Generic Receive Offload) [GRO+GSO](#)
- GSO (Generic Segmentation Offload) [GRO+GSO with DPDK](#)
- Use of multiple hardware queues [Multiqueue NIC](#), [Supplement: RSS+RPS+...](#)
- ...

Outline

- The journey of a packet through the Linux network stack
- **Need for kernel bypass techniques for packet processing**
- Kernel-bypass techniques
 - User-space packet processing
 - **Data Plane Development Kit (DPDK)**
 - **Netmap**
 - User-space network stack
 - **mTCP**
- What's trending?

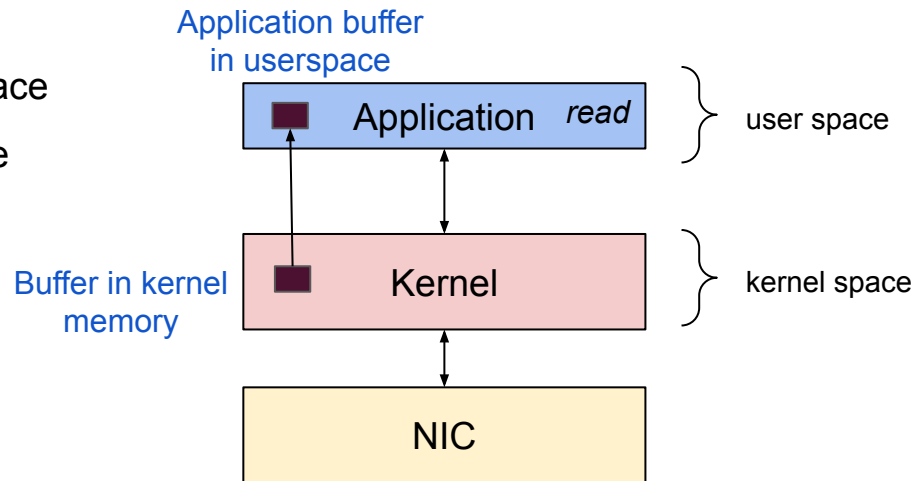
Packet Processing Overheads in Kernel

- Context switch between kernel and userspace




Packet Processing Overheads in Kernel

- Context switch between kernel and userspace
- Packet copy between kernel and userspace

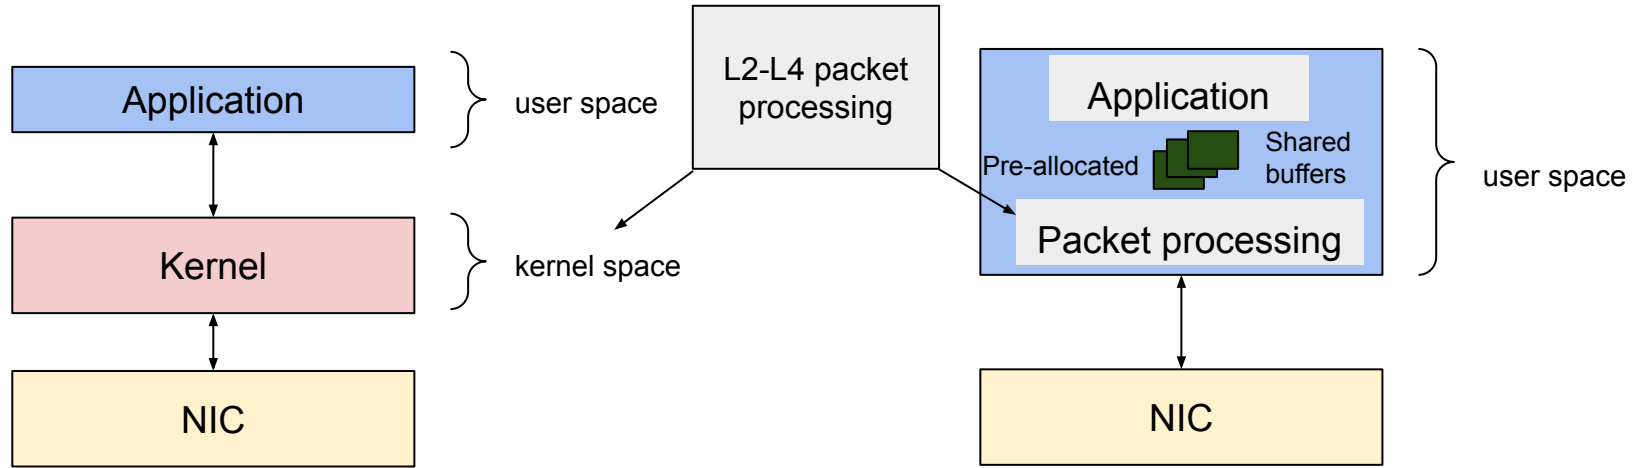


Packet Processing Overheads in Kernel

- Context switch between kernel and userspace
- Packet copy between kernel and userspace
- Dynamic allocation of sk_buff
- Per packet interrupt
- Shared data structures



Overcome Overheads in Kernel: Bypass the kernel



- ✗ Context switch between kernel and userspace
- ✗ Packet copy between kernel and userspace
- ✗ Dynamic allocation of `sk_buff`

Interrupt vs Poll Mode

Interrupt Mode



- NIC notifies it needs servicing
- Interrupt is a hardware mechanism
- Handled using interrupt handler
- Interrupt overhead for high speed traffic
- **Interrupt for a batch of packets**

Poll Mode



- CPU keeps checking the NIC
- Polling is done with help of control bits (**Command-ready** bit)
- Handled by the CPU
- Consumes CPU cycles but handles high speed traffic

Interrupt vs Poll Mode: Kernel bypass techniques

Interrupt Mode



- NIC notifies it needs servicing
- Interrupt is a hardware mechanism
- Handled using interrupt handler
- Interrupt overhead for high speed traffic

Netmap

Poll Mode



- CPU keeps checking the NIC
- Polling is done with help of control bits(**Command-ready** bit)
- Handled by the CPU
- Consumes CPU cycles but handles high speed traffic


DPDK

Outline

- The journey of a packet through the Linux network stack
- Need for kernel bypass techniques for packet processing
- **Kernel-bypass techniques**
 - User-space packet processing
 - **Data Plane Development Kit (DPDK)**
 - **Netmap**
 - User-space network stack
 - **mTCP**
- What's trending?

Intel Data Plane Development Kit (DPDK)

- Poll mode user space drivers (uio)
 - Unbinds NIC from kernel
- Mempool: HUGE pages to avoid TLB misses.
- Rte_mbuf: metadata+ pkt buffer
- Cooperative multiprocessing
 - Safe for trusted application

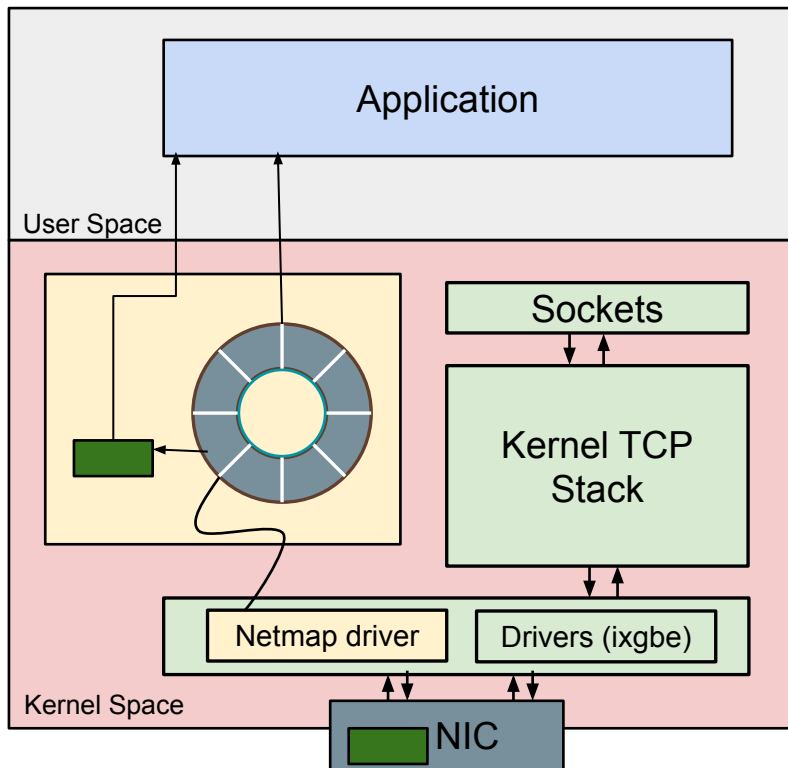


Netmap

- Netmap Rings are memory regions in kernel space shared between application and kernel
- No extra copy of a packet
- NIC can work with netmap as well as kernel drivers (transparent mode)

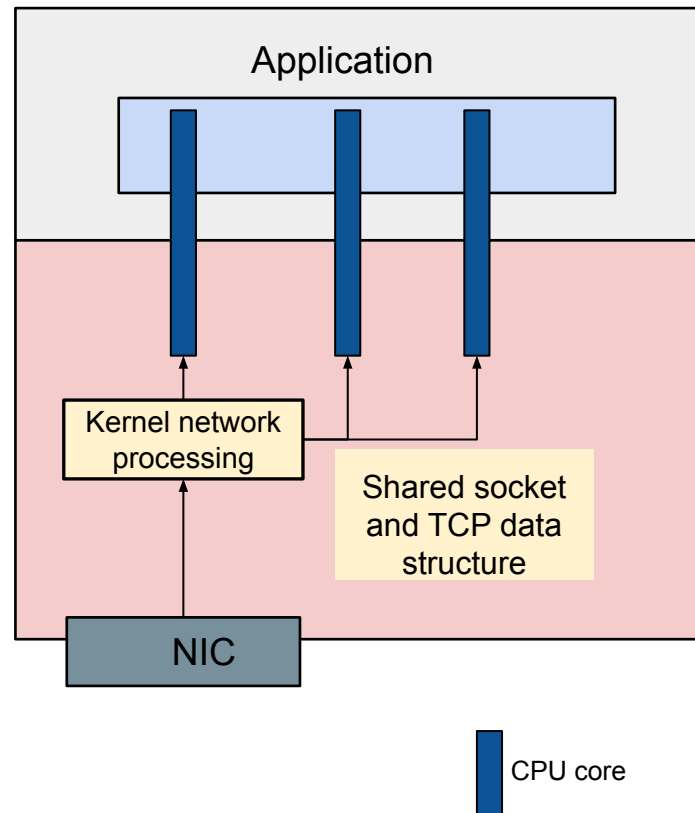
DPDK, netmap manage processing till L2 of network stack

[netmap](#)




What about L3-L7 processing?

- Overheads with L3-L7 processing in kernel
 - Shared data structure
- Userspace network stack
 - Over netmap or DPDK
- mTCP: multicore TCP



Multiqueue NIC




mTCP: Userspace network stack

- Designed for multicore scalable application
- Per core TCP data structures
 - E.g. accept queue, socket list
 - Lock free
 - Connection locality
- Leverages multiqueue support of NIC

✗ Shared data structures

mTCP



Outline

- The journey of a packet through the Linux network stack
- Need for kernel bypass techniques for packet processing
- Kernel-bypass techniques
 - User-space packet processing
 - **Data Plane Development Kit (DPDK)**
 - **Netmap**
 - User-space network stack
 - **mTCP**
- **What's trending?**

What's trending?

- Offload application processing to the kernel
 - BPF (Berkeley Packet Filter)
 - eBPF (eXtended BPF) [BPF+eBPF+XDP link-1](#), [BPF+eBPF+XDP tutorial link-2](#)
- Offload application processing to the NIC driver
 - XDP (eXpress DataPath) [Sample apps for eBPF + XDP](#)
- Offload application processing to programmable hardware
 - Programmable SmartNICs (NPU/DPU)
 - Netronome, Mellanox, Bluefield, Pensando [Video on smartNIC architecture + Netronome NIC specifics](#)
 - Programmable FPGAs
 - Xilinx, Altera
 - Programmable hardware ASICs [Programmable network: Intro video](#), [Detailed video link](#)
 - Barefoot Tofino, Cisco's Doppler, Intel Flexpipe, Cavium's Xpliant