

Monitoring Linux Network Stack

Published at 2020-04-25 | Last Update 2022-07-02

This post shows how to collect metrics from your Linux network stack (with bash scripts), and monitoring the stack status with Prometheus and Grafana.

This post assumes you have read through the following posts (kernel 3.13 + intel 1Gbps NIC driver):

1. [Monitoring and Tuning the Linux Networking Stack: Receiving Data](#)
2. [Monitoring and Tuning the Linux Networking Stack: Sending Data](#)

Or my updated versions (kernel 5.10 + mellanox 25Gbps NIC driver) if you can read Chinese:

1. [Linux 网络栈原理、监控与调优：前言](#)
2. [Linux 中断（IRQ/softirq）基础：原理及内核实现](#)
3. [Linux 网络栈接收数据（RX）：原理及内核实现](#)
4. [Linux 网络栈接收数据（RX）：配置调优](#)

Besides, some basic understandings of Prometheus and Grafana are needed.

- 1. NIC
 - 1.1 Data sources
 - `ethtool -S`
 - `/proc/net/dev`
 - `/sys/class/net/<nic>/statistics`
 - 1.2 Metrics
 - 1.3 Panels
- 2. Hardware Interrupts
 - 2.1 Data source (`/proc/interrupts`)
 - 2.2 Metric

- 2.3 Panels
- 3. Software Interrupts (softirq)
 - 3.1 Data source (/proc/softirqs)
 - 3.2 Metric
 - 3.3 Panels
- 4. Kernel Processing Drops
 - 4.1 Data source (/proc/net/softnet_stat)
 - 4.2 Metric
 - 4.3 Panel
- 5. L3 statistics (IPv4)
 - 5.1 Data source (/proc/net/snmp)
- 6. L4 statistics (TCP)
 - 6.1 Data source
 - 6.2 Metric
 - 6.3 Panel
- 7. L4 statistics (UDP)
 - 7.1 Data source
 - /proc/net/snmp
 - /proc/net/udp
- 8. Top N nodes
- 9. More metrics
- Appendix

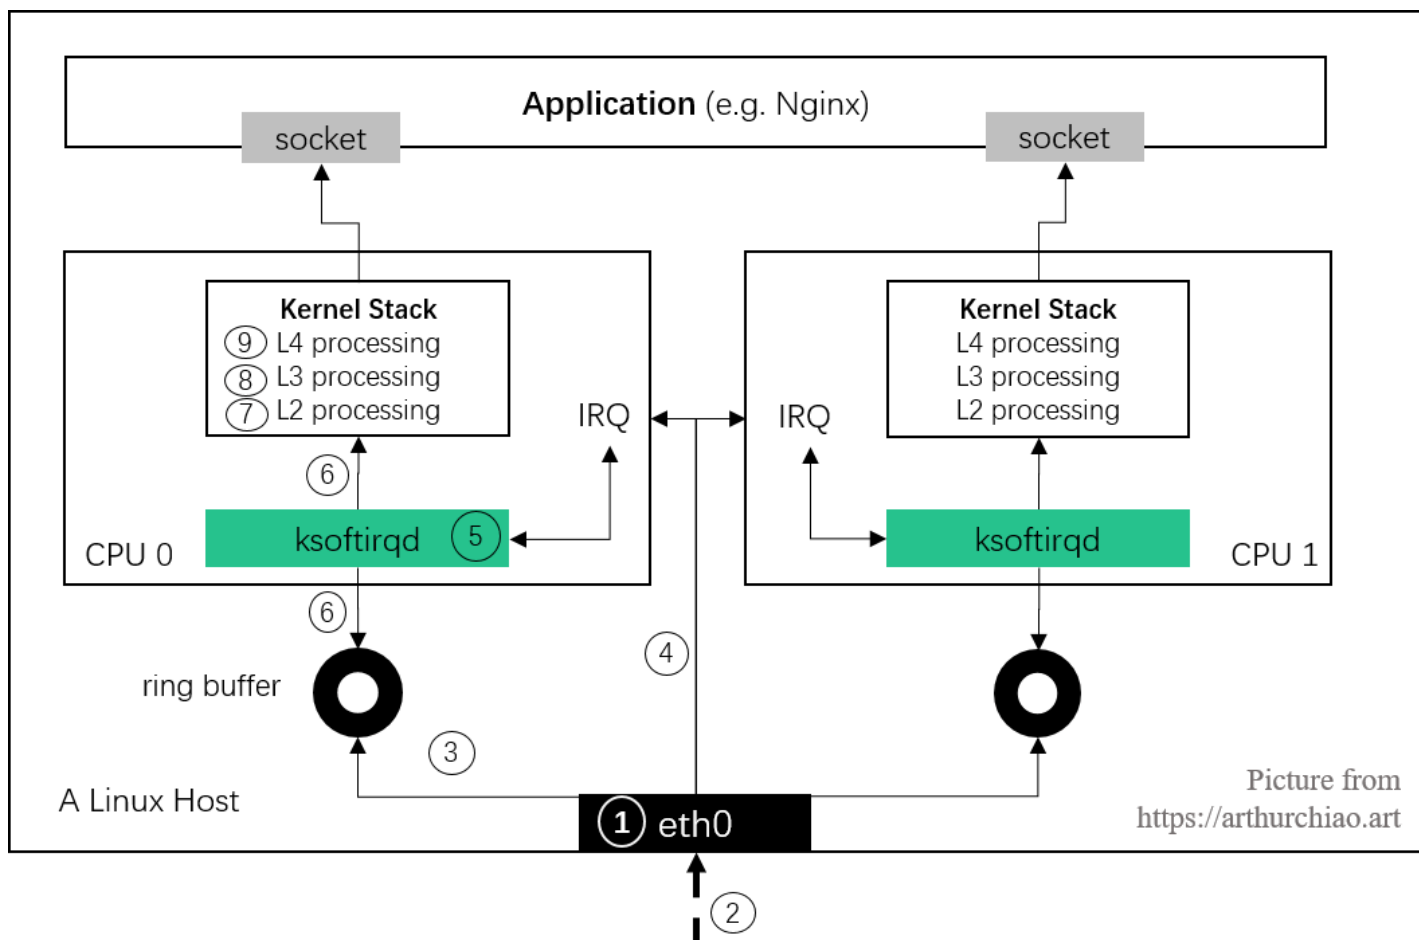


Fig. Steps of Linux kernel receiving data process and the corresponding chapters in this post

1. NIC

1.1 Data sources

ethtool -S

```
$ sudo ethtool -S eth0
NIC statistics:
  rx_packets: 597028087
  tx_packets: 5924278060
  rx_bytes: 112643393747
  tx_bytes: 990080156714
  rx_broadcast: 96
  tx_broadcast: 116
  rx_multicast: 20294528
  ....
```

Note that depending on specific NICs/drivers, some of the metrics may be updated in the kernel code (software), some may be in hardware. Refer to the datasheets of the your NICs and the corresponding drivers if needed.

/proc/net/dev

`/proc/net/dev` also provides some network device level statistics:

```
$ cat /proc/net/dev
```

Inter-	Receive									Transmit				
face	bytes	packets	errs	drop	fifo	frame	compressed	multicast	bytes	packets	errs	drop	fi	
eth0:	152214	9700	0	2	0	0	0	203860	9905984	62604	0	0	0	
lo:	463836	57535	0	0	0	0	0	0	4263836	18535	0	0	0	

Actually this is only a subnet of `/sys/class/net/<nic>/statistics` , as introduced in the below section.

/sys/class/net/<nic>/statistics

Statistics from `sysfs` provides an upper layer view of the NIC RX/TX stats.

NIC stats from `/sys/class/net/<nic>/statistics` :

```
$ ls /sys/class/net/eth0/statistics
```

collisions	rx_dropped	rx_missed_errors	tx_bytes	tx_fifo_errors
multicast	rx_errors	rx_nohandler	tx_carrier_errors	tx_heartbeat_errors
rx_bytes	rx_fifo_errors	rx_over_errors	tx_compressed	tx_packets
rx_compressed	rx_frame_errors	rx_packets	tx_dropped	tx_window_errors
rx_crc_errors	rx_length_errors	tx_aborted_errors	tx_errors	

```
$ cat /sys/class/net/eth0/statistics/rx_crc_errors
0
```

Again, look into the specific driver code if you'd like to ensure when and where a specific counter is updated.

1.2 Metrics

We will arrange our metrics in **Prometheus format**:

```
$ cat collect-network-stats.sh
#!/bin/bash

PREFIX="network"

nic_stats_output() {
    NIC=$1
    METRIC=$PREFIX"_nic_stats";

    for f in $(ls /sys/class/net/$NIC/statistics/); do
        TAGS="{\"nic\":\"$NIC\", \"type\":\"$f\"}";
        VAL=$(cat /sys/class/net/$NIC/statistics/$f 2>/dev/null);
        echo $METRIC$TAGS $VAL;
    done
}


nic_stats_output eth0
nic_stats_output eth1
```

Test:

```
$ ./collect-network-stats.sh
network_nic_stats{"nic":"eth0", "type":"collisions"} 0
network_nic_stats{"nic":"eth0", "type":"multicast"} 17775912
network_nic_stats{"nic":"eth0", "type":"rx_bytes"} 322700688616
network_nic_stats{"nic":"eth0", "type":"rx_compressed"} 0
network_nic_stats{"nic":"eth0", "type":"rx_crc_errors"} 0
network_nic_stats{"nic":"eth0", "type":"rx_dropped"} 0
network_nic_stats{"nic":"eth0", "type":"rx_errors"} 0
...
```

1.3 Panels

Push the metrics into your prometheus server, or, configure your prometheus to pull this data, we could create a panel like this:



Where, the Grafane query is:

```
avg(irate(network_nic_stats{host="$hostname"})*60) by (nic, type)
```

Note that our collecting agent automatically added a new tag `host=<hostname>` to all the metrics, so we could filter the metrics with `host="$hostname"` .

2. Hardware Interrupts

2.1 Data source (/proc/interrupts)

Information included:

- 1. IRQ->CPU mapping
- 2. Total count of specific interrupts (e.g. RX/TX interrupts, timer interrupts), on specific CPU
- 3. IRQ handler type and name

Note that these are interrupt counts, not number of packets - for example, if Interrupt Coalescing is enabled, a batch of packets will trigger only one IRQ.

```
$ cat /proc/interrupts
          CPU0      CPU1      CPU2      ...      CPU30      CPU31
...
139:         0         0         0      ...         0         0  IR-PCI-MSI 1572864-edge  eth
140:         0         0         0      ...         0         0  IR-PCI-MSI 1572865-edge  eth
141:         0         0         0      ...         0         0  IR-PCI-MSI 1572866-edge  eth
142:         0         0      308405  ...         0         0  IR-PCI-MSI 1572867-edge  eth
143:         0         0         0      ...         0         0  IR-PCI-MSI 1572868-edge  eth
```

144:	0	0	4	...	0	0	IR-PCI-MSI	1574912-edge	eth
145:	0	0	0	...	0	0	IR-PCI-MSI	1574913-edge	eth
146:	0	38	673	...	0	0	IR-PCI-MSI	1574914-edge	eth
147:	0	75604	0	...	0	0	IR-PCI-MSI	1574915-edge	eth
148:	0	3086	824199	...	0	0	IR-PCI-MSI	1574916-edge	eth
...									

2.2 Metric

Add this code snippet to our script:

```
interrupts_output() {
    PATTERN=$1
    METRIC=$PREFIX"_interrupts_by_cpu"

    egrep "$PATTERN" /proc/interrupts | awk -v metric=$METRIC \
        '{ for (i=2;i<=NF-3;i++) sum[i]+=$i;}
        END {
            for (i=2;i<=NF-3; i++) {
                tags=sprintf("{\"cpu\":\"%d\"}", i-2);
                printf(metric tags " " sum[i] "\n");
            }
        }'

    METRIC=$PREFIX"_interrupts_by_queue"
    egrep "$PATTERN" /proc/interrupts | awk -v metric=$METRIC \
        '{ for (i=2;i<=NF-3; i++)
            sum+=$i;
            tags=sprintf("{\"queue\":\"%s\"}", $NF);
            printf(metric tags " " sum "\n");
            sum=0;
        }'
}

# interface patterns
# eth: intel
# mlx: mellanox
interrupts_output "eth|mlx"

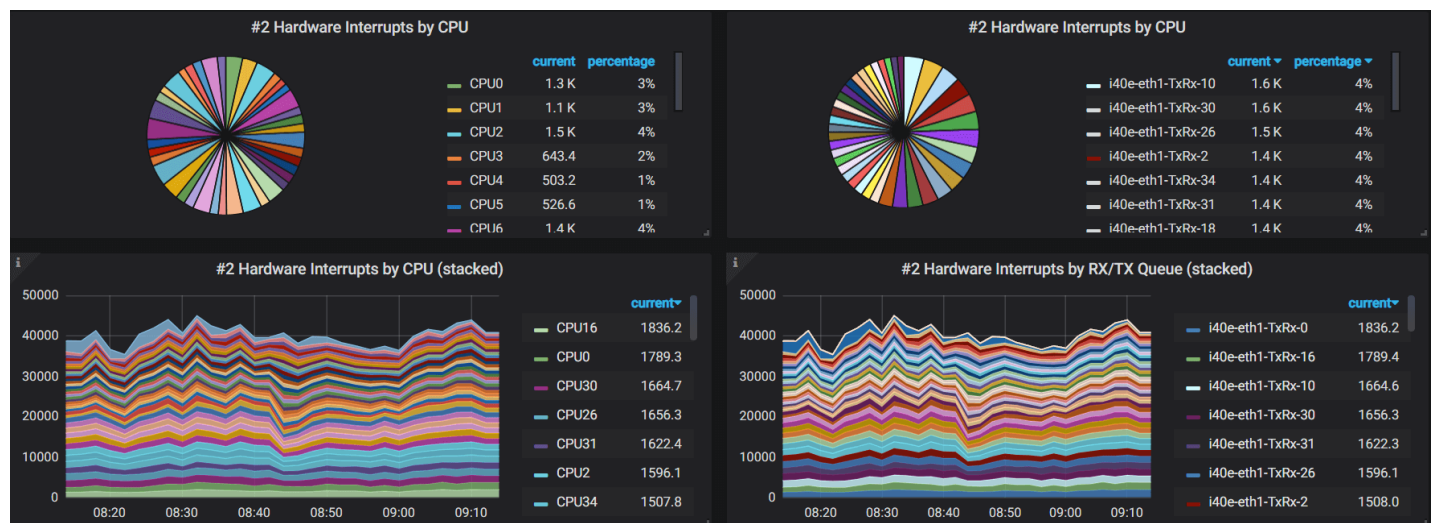
$ ./collect-network-stats.sh
network_interrupts_by_cpu{"cpu":"0"} 0
network_interrupts_by_cpu{"cpu":"1"} 6078192
network_interrupts_by_cpu{"cpu":"2"} 85118785
...
```

```

network_interrupts_by_queue{"queue":"eth0-tx-0"} 190533384
network_interrupts_by_queue{"queue":"eth0-rx-1"} 26873848
network_interrupts_by_queue{"queue":"eth0-rx-2"} 23715431
network_interrupts_by_queue{"queue":"eth0-rx-3"} 87702361
...
network_interrupts_by_queue{"queue":"eth1-rx-4"} 3119407

```

2.3 Panels



```
avg(irate(network_interrupts_by_cpu{host=~"$hostname"})) by (cpu)
```

```
avg(irate(network_interrupts_by_queue{host=~"$hostname"})) by (queue)
```

3. Software Interrupts (softirq)

3.1 Data source (/proc/softirqs)

```
$ cat /proc/softirqs
```

	CPU0	CPU1	...	CPU62	CPU63
HI:	1	0	...	0	0
TIMER:	20378862	2149097	...	0	0
NET_TX:	5	1	...	0	0
NET_RX:	1179	1868	...	0	0
BLOCK:	88034	33007	...	0	0
IRQ_POLL:	0	0	...	0	0
TASKLET:	22	0	...	0	0

SCHED:	13906041	1474443	...	0	0
HRTIMER:	0	0	...	0	0
RCU:	12121418	1964562	...	0	0

3.2 Metric

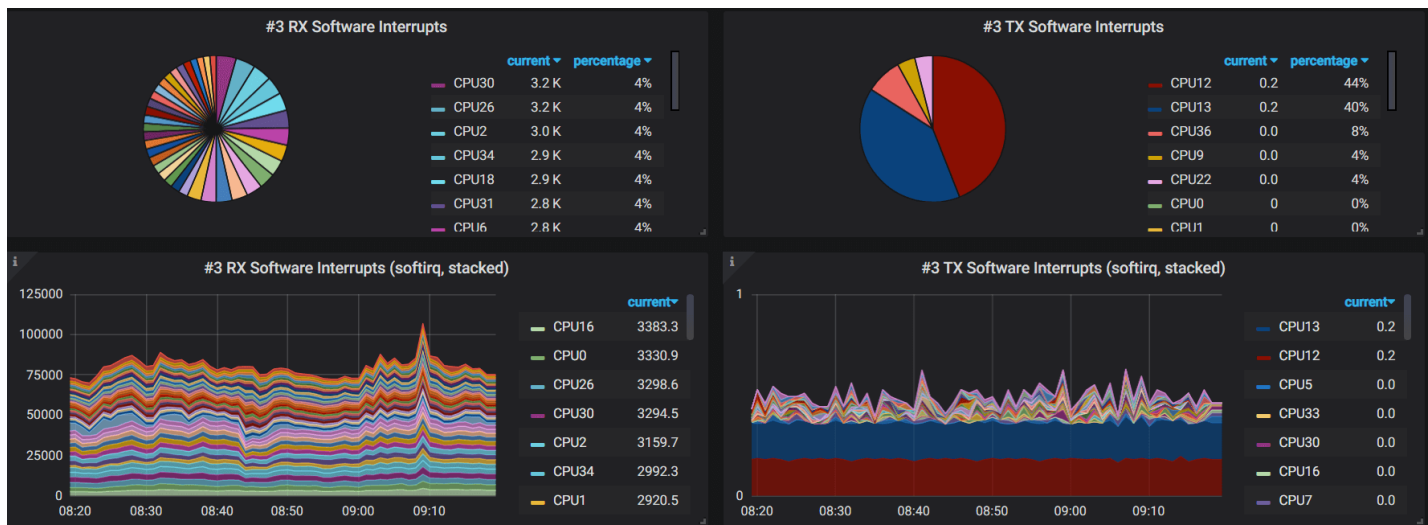
```
softirqs_output() {
    METRIC=$PREFIX"_softirqs"

    for dir in "NET_RX" "NET_TX"; do
        grep $dir /proc/softirqs | awk -v metric=$METRIC -v dir=$dir \
            '{ for (i=2;i<=NF-1;i++) {
                tags=sprintf("{\"cpu\": \"%d\", \"direction\": \"%s\"}", i-2, dir); \
                printf(metric tags " " $i "\n"); \
            }
        }'
    done
}

softirqs_output
```

```
$ ./collect-network-stats.sh
network_softirqs{"cpu": "0", "direction": "NET_RX"} 196082
network_softirqs{"cpu": "1", "direction": "NET_RX"} 119888284
network_softirqs{"cpu": "2", "direction": "NET_RX"} 189840914
network_softirqs{"cpu": "3", "direction": "NET_RX"} 114621858
network_softirqs{"cpu": "4", "direction": "NET_RX"} 1453599
network_softirqs{"cpu": "5", "direction": "NET_RX"} 192694791
network_softirqs{"cpu": "6", "direction": "NET_RX"} 49328487
...
```

3.3 Panels



Grafana queries:

```
avg(irate(network_interrupts_by_cpu{host="$hostname",direction="NET_RX"})) by (cpu)
```

```
avg(irate(network_interrupts_by_cpu{host="$hostname",direction="NET_TX"})) by (cpu)
```

4. Kernel Processing Drops

4.1 Data source (/proc/net/softnet_stat)

```
$ cat /proc/net/softnet_stat
```

```
00049007 00000000 00000000 00000000 00000000 00000000 00000000 00000000 00000000 00000000 000
074c3e6e 00000000 00000000 00000000 00000000 00000000 00000000 00000000 00000000 00000000 000
0c98d81b 00000000 00000000 00000000 00000000 00000000 00000000 00000000 00000000 00000000 000
07212d42 00000000 00000000 00000000 00000000 00000000 00000000 00000000 00000000 00000000 000
0018ad7c 00000000 00000000 00000000 00000000 00000000 00000000 00000000 00000000 00000000 000
00037314 00000000 00000002 00000000 00000000 00000000 00000000 00000000 00000000 00000000 000
...
```

4.2 Metric

```
softnet_stat_output() {
    TYP=$1
    IDX=$2
```

```

METRIC=$PREFIX"_softnet_stat"

VAL=$(cat /proc/net/softnet_stat | awk -v IDX="$IDX" '{sum+=strtonum("0x"$IDX);} END{prin
TAGS="{\"type\": \"$TYP\"}";

echo $METRIC$TAGS $VAL;
}

# Format of /proc/net/softnet_stat:
#
# column 1 : received frames
# column 2 : dropped
# column 3 : time_squeeze
# column 4-8: all zeros
# column 9 : cpu_collision
# column 10 : received_rps
# column 11 : flow_limit_count
#
# http://arthurchiao.art/blog/tuning-stack-rx-zh/
softnet_stat_output "dropped" 2
softnet_stat_output "time_squeeze" 3
softnet_stat_output "cpu_collision" 9
softnet_stat_output "received_rps" 10
softnet_stat_output "flow_limit_count" 11

```


Run:

```

$ ./collect-network-stats.sh
network_softnet_stat{"type":"dropped"} 0
network_softnet_stat{"type":"time_squeeze"} 4
network_softnet_stat{"type":"cpu_collision"} 0
network_softnet_stat{"type":"received_rps"} 0
network_softnet_stat{"type":"flow_limit_count"} 0

```

4.3 Panel



Grafana queries:


```
avg(irate(network_softnet_stat{host="$hostname"})) by (type)
```

5. L3 statistics (IPv4)

5.1 Data source (/proc/net/snmp)

```
$ cat /proc/net/snmp
```

```
Ip: Forwarding DefaultTTL InReceives InHdrErrors InAddrErrors ForwDatagrams InUnknownProtos I
Ip: 1 64 25922988125 0 0 15771700 0 0 25898327616 22789396404 12987882 51 1 10129840 2196520
...
```



Fields:

```
// https://github.com/torvalds/linux/blob/v5.10/include/uapi/linux/snmp.h#L21
```

```
/* ipstats mib definitions */
```

```
enum {
    IPSTATS_MIB_NUM = 0,
    IPSTATS_MIB_INPKTS,      /* InReceives */
    IPSTATS_MIB_INOCTETS,    /* InOctets */
    IPSTATS_MIB_INDELIVERS,  /* InDelivers */
    ...
}
```

6. L4 statistics (TCP)

6.1 Data source

```
$ netstat -s
Ip:
  397147220 total packets received
  621 with invalid headers
  1 with invalid addresses
  16591642 forwarded
...
Tcp:
  53687405 active connections openings
  449771 passive connection openings
  52888864 failed connection attempts
  66565 connection resets received
...
TcpExt:
  18 ICMP packets dropped because they were out-of-window
  4 ICMP packets dropped because socket was locked
  643745 TCP sockets finished time wait in fast timer
  8 packets rejects in established connections because of timestamp
...
```

6.2 Metric

```
netstat_output() {
  PATTERN=$1
  ARG_IDX=$2

  METRIC=$PREFIX"_tcp"
  VAL=$(netstat -s | grep "$PATTERN" | awk -v i=$ARG_IDX '{print $i}')

  # generate "type" string with prefix and pattern
  #
  # 1. replace whitespaces with underlines
  # 2. remove trailing dollar symbol ('$') if there is
  #
  # e.g. "fast retransmits$" -> "fast_retransmits"
  #
  TYP=$(echo "$PATTERN" | tr ' ' '_' | sed 's/\$/g')

  TAGS="{\"type\": \"$TYP\"}";
  echo $METRIC$TAGS $VAL;
}
```

```

netstat_output "segments retransmitted" 1
netstat_output "TCPLostRetransmit" 2
netstat_output "fast retransmits$" 1
netstat_output "retransmits in slow start" 1
netstat_output "classic Reno fast retransmits failed" 1
netstat_output "TCPSynRetrans" 2

netstat_output "bad segments received" 1
netstat_output "resets sent$" 1
netstat_output "connection resets received$" 1

netstat_output "connections reset due to unexpected data$" 1
netstat_output "connections reset due to early user close$" 1

```


Run:

```

$ ./collect-network-stats.sh
network_tcp{"type":"segments_retransmitted"} 618183
network_tcp{"type":"TCPLostRetransmit"} 133668
network_tcp{"type":"fast_retransmits"} 45745
network_tcp{"type":"retransmits_in_slow_start"} 62977
network_tcp{"type":"classic_Reno_fast_retransmits_failed"} 418
network_tcp{"type":"TCPSynRetrans"} 175919
network_tcp{"type":"bad_segments_received"} 399
network_tcp{"type":"resets_sent"} 234094
network_tcp{"type":"connection_resets_received"} 66553
network_tcp{"type":"connections_reset_due_to_unexpected_data"} 93589
network_tcp{"type":"connections_reset_due_to_early_user_close"} 6522

```

6.3 Panel



Grafana queries:

```
avg(irate(network_tcp{host="$hostname"})*60) by (type)
```

7. L4 statistics (UDP)

7.1 Data source

/proc/net/snmp

```
$ cat /proc/net/snmp
```

```
...
```

```
Udp: InDatagrams NoPorts InErrors OutDatagrams RcvbufErrors SndbufErrors InCsumErrors Ignored
```

```
Udp: 3251496 356 0 3251774 0 0 0 0
```

```
UdpLite: InDatagrams NoPorts InErrors OutDatagrams RcvbufErrors SndbufErrors InCsumErrors Ign
```

```
UdpLite: 0 0 0 0 0 0 0 0
```

/proc/net/udp

```
$ cat /proc/net/udp
```

s1	local_address	rem_address	st	tx_queue	rx_queue	tr	tm->when	retrnsmt	uid	timeout	in
52544:	3F40160A:007B	00000000:0000	07	00000000:00000000	00:00000000	00000000	00000000	00000000	38	0	1
52544:	016C000A:007B	00000000:0000	07	00000000:00000000	00:00000000	00000000	00000000	00000000	38	0	8
52544:	7136060A:007B	00000000:0000	07	00000000:00000000	00:00000000	00000000	00000000	00000000	38	0	5
52544:	0100007F:007B	00000000:0000	07	00000000:00000000	00:00000000	00000000	00000000	00000000	0	0	5
52544:	00000000:007B	00000000:0000	07	00000000:00000000	00:00000000	00000000	00000000	00000000	0	0	5

The first line describes each of the fields in the lines following:

- `s1` : Kernel hash slot for the socket
- `local_address` : Hexadecimal local address of the socket and port number, separated by `:`.
- `rem_address` : Hexadecimal remote address of the socket and port number, separated by `:`.
- `st` : The state of the socket. Oddly enough, the UDP protocol layer seems to use some TCP socket states. In the example above, 7 is `TCP_CLOSE`.
- `tx_queue` : The amount of memory allocated in the kernel for outgoing UDP datagrams.
- `rx_queue` : The amount of memory allocated in the kernel for incoming UDP datagrams.

- `tr`, `tm->`when, `retrnsmt`: These fields are unused by the UDP protocol layer.
- `uid`: The effective user id of the user who created this socket.
- `timeout`: Unused by the UDP protocol layer.
- `inode`: The inode number corresponding to this socket. You can use this to help you determine which user process has this socket open. Check `/proc/[pid]/fd`, which will contain symlinks to `socket[:inode]`.
- `ref`: The current reference count for the socket.
- `pointer`: The memory address in the kernel of the struct sock.
- `drops`: The number of datagram drops associated with this socket. Note that this does not include any drops related to sending datagrams (on corked UDP sockets or otherwise); this is only incremented in receive paths as of the kernel version examined by this blog post.

Printing code: `net/ipv4/udp.c`.

8. Top N nodes

Top-N nodes for some specific metrics:



This is quite helpful for detecting problematic nodes.

Queries for these panels are very similar, we list some here:


- `topk(10, avg(irate(k8s.node.network.tcp{type="segments_retransmitted"})*60) by (host))`

- `topk(10, avg(irate(k8s.node.network.tcp{type="TCPSynRetrans"})*60) by (host))`
- `topk(10, avg(irate(k8s.node.network.nic.errors{type=~'rx_.*'})*60) by (host))`
- `topk(10, avg(irate(k8s.node.network.nic.errors{type=~'tx_.*'})*60) by (host))`

9. More metrics

This post serves as a introductory guide for how to monitoring you network stack with Prometheus and Grafana.

Actually there are more metrics than we have shown in the above, such as, you could monitor the NIC bandwidth with metrics from NIC statistics:



Besides, you could also configure alerting rules on Grafana panels, e.g. alerting when NIC errors exceeds a pre defined threshold.

Appendix

1. [collect-network-stats.sh](#)

« **L4LB FOR KUBERNETES: THEORY AND PRACTICE WITH CILUM+BGP+ECMP**

© 2016-2022 [Arthur Chiao](#), Powered by [Jekyll](#), Theme originated from [Long Haul](#). Site visits: 631915, powered by [busuanzi](#)

