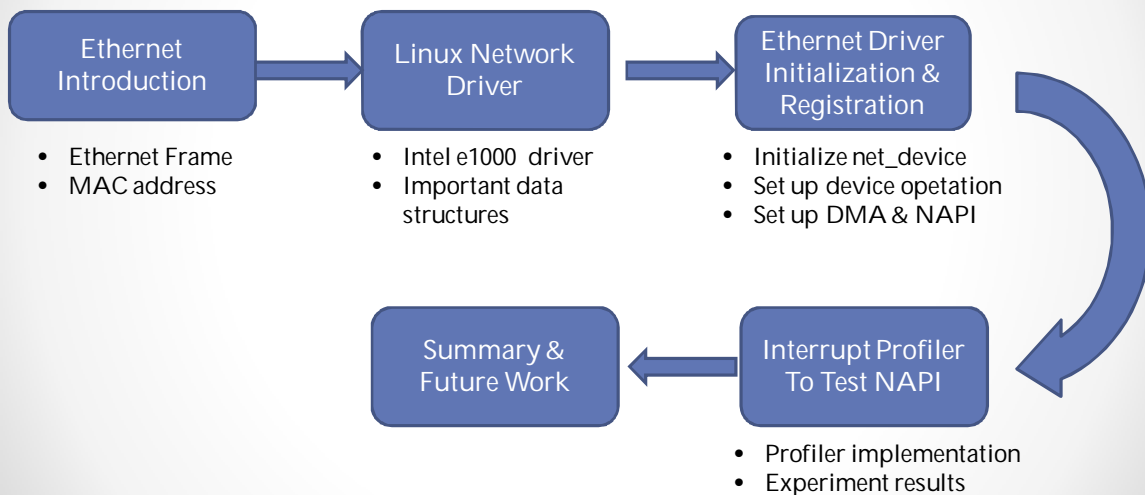


Understanding Linux Network Device Driver and NAPI Mechanism

Xinying Wang, Cong Xu
CS 423 Project

Outline




Introduction to Ethernet

- A family of computer networking technologies for local area networks (LANs)
- Commercially introduced in 1980 and standardized in 1983 as IEEE 802.3.
- The most popular network with good degree of compatibility
- Features:
 - Ethernet frame
 - MAC Address



Ethernet frame

- Transported by Ethernet packet (a data packet on an Ethernet)
- Example of Ethernet frame structure through TCP socket:




- Ethernet header
 - Header: a set of bytes (octets*) prepended to a packet
 - Include destination MAC address and source MAC address
- FCS: to detect any in-transit corruption of data

*octet: a group of eight bits

MAC address

- Media Access Control address
- Often stored in hardware's read-only memory
- First three octets: Organizationally Unique Identifier (OUI)
- Following octets: as long as unique



Linux network driver

- Linux kernel handles MAC address resolution.
- Network drivers are still needed
 - Kernel cannot do anything
 - Different from character drivers and block drivers
- Intel e1000 driver for Ethernet adapter
 - `/drivers/net/ethernet/intel/e1000`



Data structure: *struct net_device*

- Global information
 - char name[IFNAMSIZ]:
 - The name of the device.
 - unsigned long state:
 - Device state.
 - struct net_device *next:
 - Pointer to the next device in the global linked list.
 - int (*init)(struct net_device *dev):
 - An initialization function.

Data structure: *struct net_device*

- Hardware information:
 - unsigned long rmem_end, rmem_start, mem_end, mem_start:
 - Device memory information.
 - unsigned long base_addr:
 - The I/O base address of the network interface.
 - unsigned char irq:
 - The assigned interrupt number.
 - unsigned char dma:
 - The DMA channel allocated by the device.

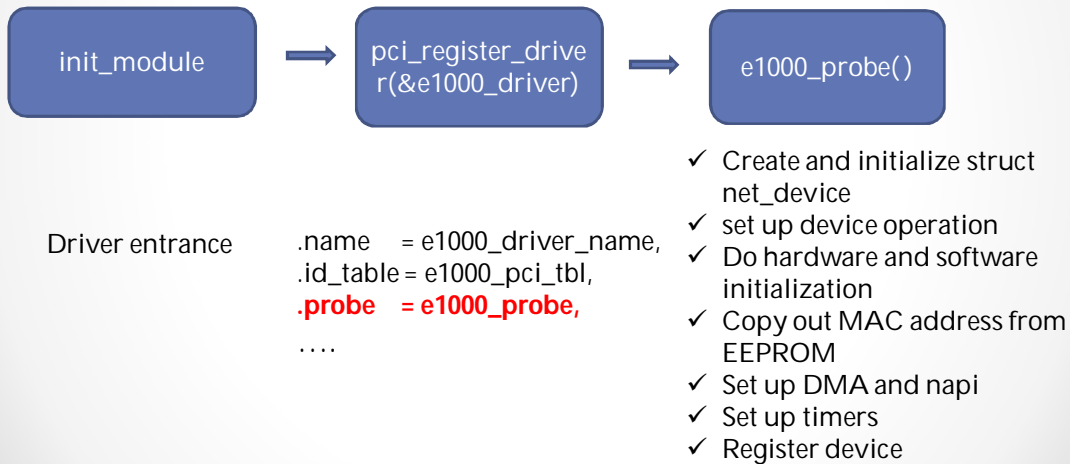
Data structure: *struct e1000_adapter*

- `struct net_device *netdev;`
 - Pointer to `net_device` struct
- `struct pci_dev *pdev;`
 - Pointer to `pci_device` struct
- `struct e1000_hw hw;`
 - An `e1000_hw` struct
- `struct e1000_hw_stats stats;`
 - Statistics counters collected by the MAC

Data structure: *struct e1000_hw*

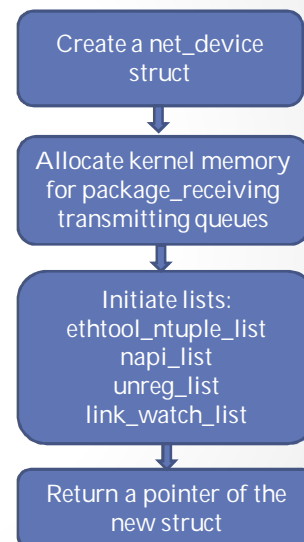
- `e1000_mac_type mac_type;`
 - An enum for currently available devices
- `u8 mac_addr[NODE_ADDRESS_SIZE];`
 - MAC address
- `u16 device_id;`
 - Device identification information
- `u16 vendor_id;`
 - Vendor information

Ethernet driver initialization and registration



Initialize struct net_device

- Initialization is done by calling MACRO
alloc_etherdev(sizeof_priv)
- Track down to function
struct net_device
*alloc_netdev_mq(int sizeof_priv,
const char *name, void
(*setup)(struct net_device *),
unsigned int queue_count) in
net/core/dev.c
- What does this function do?



Set up device operation

- It is defined in struct net_device_ops
- What does device operation do?
 - open
 - Close
 - Get System Network Statistics
 - Configuring hardware for uni or multicast
 - Change Ethernet address
 - Set transmission time-out
 - Change MTU
 - I/O control
 - Validate Ethernet address

Hardware and software initialization

- Hardware initialization

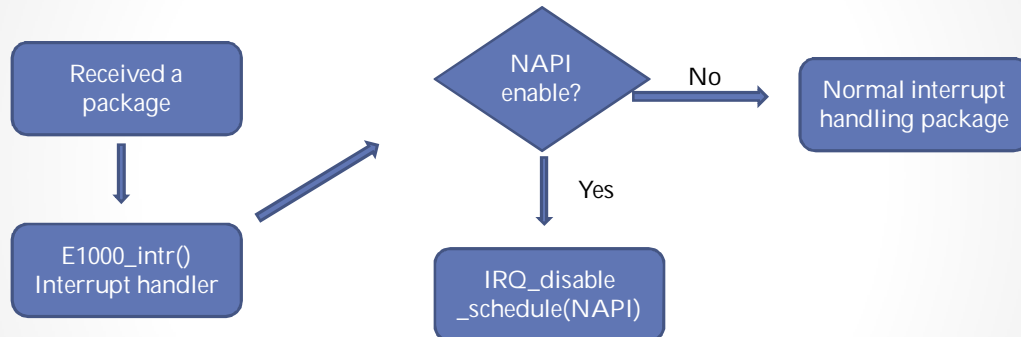
Initialize members of hw struct; abstract vendor ID, device ID, subsystem ID; identify mac type; set MTU size.
- Software initialization

This is done after hardware initialization; Initialize general software structures (struct e1000_adapter)

Set up DMA and NAPI

- What is NAPI and why do we need NAPI?
- Allocate buffer skb
e1000_rx_ring
- Remap DMA
dma_map_single()
- NAPI add
netif_napi_add(struct net_device *dev, struct napi_struct *napi, int (*poll)(struct napi_struct *, int), int weight)

How a package being received



E1000 driver is NAPI-enabled

NAPI implementation

In interrupt handler function `e1000_entr()`

Make sure the net device is working properly
`netif_rx_schedule_prep()`



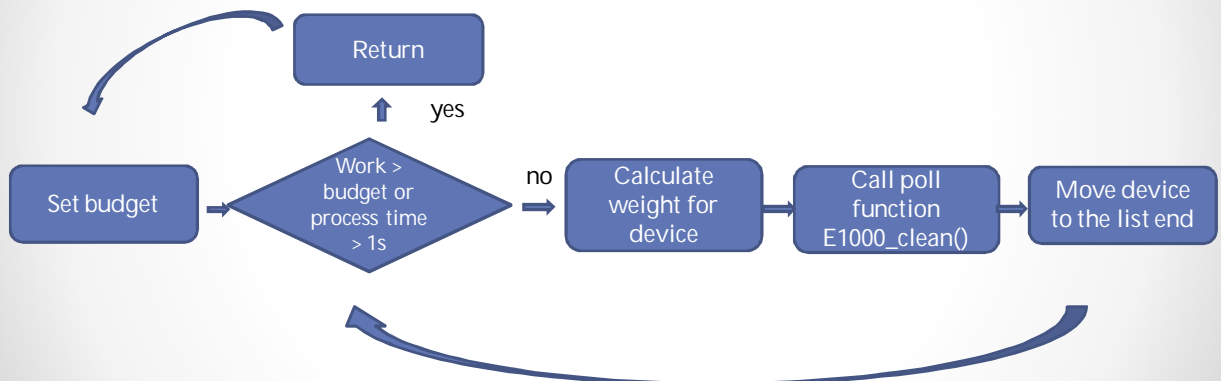
Add net device into poll list
`_netif_rx_schedule()->napi_schedule()`



`_raise_softirq_irqoff(NET_RX_SOFTIRQ)` for switching to bottom half


NAPI implementation

- Bottom half function `net_rx_action()`



NAPI implementation

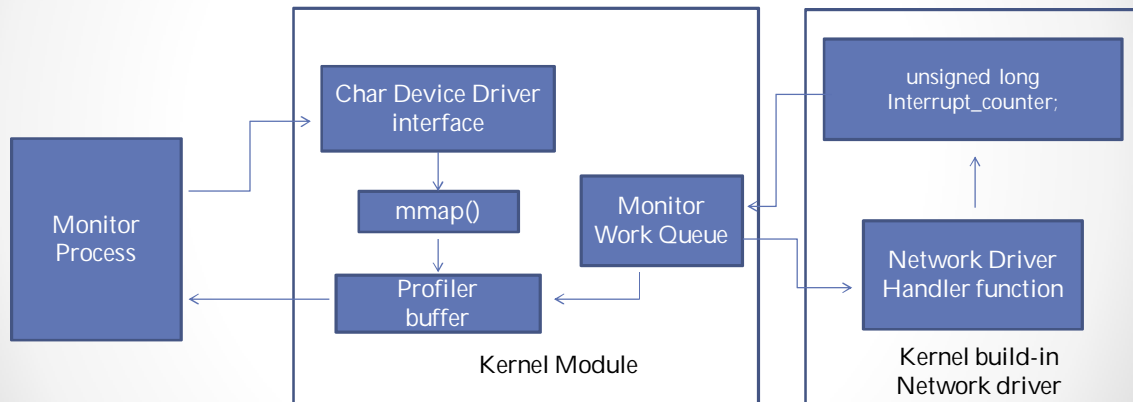
- Poll function `e1000_clean()`




Experiment

- An experiment is designed to test NAPI mechanism
- A interrupt profiler is designed to profile the interrupt counts in a designed period
- Linux kernel 3.13.6 was employed to fulfill the experiment
- Experiment platform: CPU: Intel core i5 dual core 2.53Ghz; Memory: 4G; Network card: Intel-82577 Gb card

Profiler implementation



Results



Summary and future work

- The Linux network device driver was analyzed base on Intel E1000 driver code files.
- The mechanism and implementation of NAPI was detailed
- An experiment was designed to further understand the NAPI mechanism
- A thorough understanding the Linux network device driver could be done for the future by further analysis of more sub functions.

Reference

- Branden Moore, Thomas Slabach, Lambert Schaelicke, Profiling Interrupt Handler Performance through Kernel Instrumentation, Proceedings of the 21st international conference on computer design
- Lambert Schaelicke, Al Davis, and Sally A. Mckee, Profiling I/O Interrupts in Modern Architectures
- Linux kernel source code (version 3.13.6)

Q & A

Thanks!